

MATH324 (Statistics) – Lecture Notes

McGill University

Masoud Asgharian

Winter 2019

Contents

1	Lecture 1	3
2	Lecture 2	3
3	Lecture 3	4
3.1	Markov's Inequality	4
3.2	Tchebyshev's Inequality	4
3.3	Application to Voting	7
4	Lecture 4	9
5	Lecture 5	9
6	Lecture 6	9
7	Lecture 7	9
8	Lecture 8	9
9	Lecture 9	9
10	Lecture 10	9

11 Lecture 11	9
12 Lecture 12	9

1 Lecture 0

2 Lecture 1

3 Lecture 2

3.1 Markov's Inequality

Let X be a random variable and h be a **non-negative** function; ie:

$$h : \mathbb{R} \rightarrow \mathbb{R}^+ \cup \{0\} = [0, \infty)$$

Suppose $E(h(X)) < \infty$, then for some $\lambda > 0$, we have:

$$P(h(X) \geq \lambda) \leq \frac{E[h(X)]}{\lambda} \quad (1)$$

Proof. Suppose X is a continuous random variable:

$$\begin{aligned} E[h(x)] &= \int_{\mathbb{R}} h(x) f_X(x) dx \\ &= \left(\int_{x:h(x) \geq \lambda} h(x) f_X(x) dx + \int_{x:h(x) < \lambda} h(x) f_X(x) dx \right) \\ &\geq \int_{x:h(x) \geq \lambda} h(x) f_X(x) dx && \text{since } h \geq 0 \\ &\geq \lambda \int_{x:h(x) \geq \lambda} f_X(x) dx = \lambda P(h(X) \geq \lambda) \\ \implies P(h(X) \geq \lambda) &\leq \frac{E(h(X))}{\lambda} \end{aligned}$$

The proof for the discrete case is similar. □

3.2 Tchebyshev's Inequality

Tchebyshev's Inequality is a special case of Markov's Inequality. Consider $h(x) = (x - \mu)^2$, then:

$$\begin{aligned} P(|X - \mu| \geq \lambda) &= P((X - \mu)^2 \geq \lambda^2) \\ &\leq \frac{E[(X - \mu)^2]}{\lambda^2} && \text{if } E[(X - \mu)^2] < \infty \end{aligned}$$

Let $\mu = E(X)$, then $E[(X - \mu)^2] = \text{Var}(X)$ denoted by σ_x^2 . We therefore have:

$$P(|X - \mu_x| \geq \lambda) \leq \frac{\sigma_x^2}{\lambda^2} \quad \text{where } \mu_x = E(X) \quad (2)$$

Now consider $\lambda = K\sigma_x$ where K is a known number. Then:

$$P(|X - \mu_x| \geq K\sigma_x) \geq \frac{\sigma_x^2}{K^2\sigma_x^2} = \frac{1}{K^2} \quad (3)$$

This is called **Tchbyshev's Inequality**.

Example 3.1. Suppose $K = 3$.

$$P(|X - \mu_x| \geq 3\sigma_x) \leq \frac{1}{9}$$

In other words, at least 88% of the observations are within 3 standard deviation from the population mean.

Going back to the our example:

$$X_i \sim (\mu, 1) \quad , \quad \bar{X}_n = \frac{1}{n} \sum_{i=1}^n X_i$$

We want to study $P(\epsilon \geq \delta) = P(|\bar{X}_n - \mu| \geq \delta)$, first we note that:

$$E(X_i) = \mu \quad , \quad i = 1, 2, \dots, n$$

Then:

$$\begin{aligned} E(\bar{X}_n) &= E\left(\frac{1}{n} \sum_{i=1}^n X_i\right) = \frac{1}{n} \sum_{i=1}^n E(X_i) \\ &= \frac{1}{n} \sum_{i=1}^n \mu = \frac{1}{n} (n\mu) = \frac{1}{n} \cdot n\mu \\ &= \mu \end{aligned} \quad (*)$$

Thus, using (2) we have:

$$P(|\bar{X}_n - \mu| \geq \delta) \leq \frac{Var(\bar{X}_n)}{\delta^2}$$

Now:

$$\begin{aligned} Var(\bar{X}_n) &= Var\left(\frac{1}{n} \sum_{i=1}^n X_i\right) = \frac{1}{n^2} Var\left(\sum_{i=1}^n X_i\right) \\ &= \frac{1}{n^2} \left[\sum_{i=1}^n Var(X_i) + \sum_{1 \leq i < j \leq n} \sum_{1 \leq i < j \leq n} Cov(X_i, X_j) \right] \quad \text{using Thm 5.12(b) - page 271} \\ &= \frac{1}{n^2} \sum_{i=1}^n Var(X_i) \quad \text{since } \prod_{i=1}^n X_i \\ &= \frac{1}{n^2} n Var(X) = \frac{Var(X)}{n} \quad \text{since } x_i\text{'s are identically distributed} \\ &= \frac{\delta_X^2}{n} \end{aligned} \quad (**)$$

In our case $X \sim N(\mu, 1)$ so $Var(X) = \delta_X^2 = 1$. Thus $Var(\bar{X}_n) = \frac{1}{n}$

Remark. $X \perp\!\!\!\perp Y \implies Cov(X, Y) = 0$. Note that:

$$X \perp\!\!\!\perp Y \implies E[g_1(X)g_2(Y)] = E[g_1(X)].E[g_2(Y)]$$

in particular:

$$X \perp\!\!\!\perp Y \implies E[XY] = E[X].E[Y]$$

on the other hand:

$$Cov(X, Y) = E[XY] - E(X)E(Y)$$

thus:

$$X \perp\!\!\!\perp Y \implies Cov(X, Y) = 0.$$

recall that $X \perp\!\!\!\perp Y$ means X and Y are independent, i.e. $f_{X,Y}(x, y) = f_X(x)f_Y(y)$
where $f_{X,Y}$, f_X and f_Y represent respectively the

We therefore have:

$$P(|\bar{X}_n - \mu| \geq \delta) \leq \frac{1}{n\delta^2} \quad (4)$$

Using (4) and the sample size, n , we can find an upper bound for the proportion of deviations which are greater than a given threshold δ .

We can also use (4) for Sample Size Deterministic:

Suppose δ is given and we want $P(|\bar{X}_n - \mu| \geq \delta) \leq \beta$ where β is also given. Then setting $\frac{1}{n\delta^2} = \beta$, we can estimate $n \approx \frac{1}{\beta\delta^2}$.

3.3 Application to Voting

Define $X_i = \begin{cases} 1 & \text{NDP} \\ 0 & \text{otherwise} \end{cases}$. Associated to each eligible voter in Canada we

have a binary variable X . Let $p = P(X = 1)$. So p represents the proportion of eligible voters who favor *NDP*. Of interest is often estimation of p . Suppose we have a sample of size n , X_1, X_2, \dots, X_n .

$\bar{X}_n = \frac{1}{n} \sum_{i=1}^n X_i$ is the sample proportion; The counterpart of p which can be denoted by \hat{p} . Note that:

$$\mu_x = E(X) = 1 \times P(X = 1) + 0 \times P(X = 0) = 1 - p + 0 \times (1 - p) = p$$

and:

$$E(X^2) = 1^2 \times P(X = 1) + 0^2 \times P(X = 0) = 1 - p + 0 \times (1 - p) = p$$

From (*) and (**) we find that :

$$E(\hat{p}_n) = E(\bar{X}_n)\mu_x = p$$

and:

$$\text{Var}(\hat{p}_n) = E(\bar{X}_n) = \frac{\text{Var}(X)}{n} = \frac{\sigma_X^2}{n} = \frac{p(1-p)}{n}$$

Thus using (2), we have:

$$P(|\hat{p}_n - p| \geq \delta) \leq \frac{\text{Var}(\hat{p}_n)}{\delta^2} = \frac{p(1-p)}{n\delta^2}$$

Note that the above bound on the probability of derivation depends on p which is *unknown*. We however notice that $p(1-p) \leq \frac{1}{4}$.

Define $\mathcal{C}(x) = x(1-x)$ for $0 < x < 1$. Then:

$$\begin{aligned} \mathcal{C}'(x) &= 1 - 2x \implies \mathcal{C}'(x) = 0 \implies x = \frac{1}{2} \\ \mathcal{C}''\left(\frac{1}{2}\right) &= -2 \implies x = \frac{1}{2} \quad \text{which is a **maximizer**} \\ \mathcal{C}\left(\frac{1}{2}\right) &= \frac{1}{2}\left(1 - \frac{1}{2}\right) = \frac{1}{4} \end{aligned}$$

(Note that $\mathcal{C}''(x) = -2$ for all $0 < x < 1$)

We therefore find:

$$P(|\hat{p}_n - p| \geq \delta) \leq \frac{1}{4n\delta^2} \tag{5}$$

Using (5) and a given sample size n we can find an upper bound for the probability of derivation by δ and the amount for any given δ .

We can also use (5) for sample size deterministic for a size bound β and derivative δ as follows:

$$\frac{1}{4n\delta^2} = \beta \implies n \geq \frac{1}{4\beta\delta^2}$$

This is of course conservative since $p(1-p) \leq \frac{1}{4}$.

4 Lecture 3

4.1 MSE

MSE: To study estimation error we started by studying $P(|\hat{\Theta}_n - \Theta| > \delta)$, deviation above a given threshold δ , by bounding this probability. One may take a different approach by studying average Euclidean distance, i.e. $E[|\hat{\Theta}_n - \Theta|^2]$, which denoted by $\text{MSE}(\hat{\Theta}_n)$.

We note that if $\Theta = E(\hat{\Theta}_n)$, i.e. $\hat{\Theta}_n$ is an unbiased estimation of Θ , then:

$$\text{MSE}(\hat{\Theta}_n) = E[|\hat{\Theta}_n - \Theta|^2] = E[(\hat{\Theta}_n - \mu_{\Theta_n})^2] = \text{Var}(\hat{\Theta}_n)$$

Now recall that $\text{Var}(X) = 0 \implies P(X = \text{constant}) = 1$ which essentially means random variable X is a constant.

The same comment applies to $\text{MSE}(\hat{\Theta}_n)$. We want to find the closest estimator $\hat{\Theta}_n$ to Θ which means that we want to minimize $E[(\hat{\Theta}_n - \Theta)^2]$ over all possible estimators, ideally at least the above comment tells us that in real applications we cannot expect to find an estimator whose MSE is equal to zero. Let's try to understand the MSE a bit more:

$$\begin{aligned} \text{MSE}(\hat{\Theta}_n) &= E[(\hat{\Theta}_n - \Theta)^2] \\ &= E\left[\left((\hat{\Theta}_n - E(\hat{\Theta}_n)) + (E(\hat{\Theta}_n) - \Theta)\right)^2\right] \\ &= E\left[(\hat{\Theta}_n - E(\hat{\Theta}_n))^2\right] + \underbrace{E[(E(\hat{\Theta}_n) - \Theta)^2]}_{\text{not a r.v.}} + 2 \cdot \underbrace{E[(\hat{\Theta}_n - E(\hat{\Theta}_n)) \cdot (E(\hat{\Theta}_n) - \Theta)]}_{\text{not a r.v.}} \\ &= E[(\hat{\Theta}_n - E(\hat{\Theta}_n))^2] + E[(E(\hat{\Theta}_n) - \Theta)^2] + 2 \cdot E[(E(\hat{\Theta}_n) - \Theta) \cdot (\hat{\Theta}_n - E(\hat{\Theta}_n))] \\ &= \underbrace{\text{Var}(\hat{\Theta}_n)}_{\text{Bias}(\hat{\Theta}_n)} + \underbrace{[E(\hat{\Theta}_n) - \Theta]^2}_{\text{Bias}(\hat{\Theta}_n)} + 2 \cdot \underbrace{\text{Bias}(\hat{\Theta}_n) \cdot E[(\hat{\Theta}_n - E(\hat{\Theta}_n))]}_{E(\hat{\Theta}_n) - E(\hat{\Theta}_n) = 0} \\ &= \text{Var}(\hat{\Theta}_n) + \text{Bias}^2(\hat{\Theta}_n) \end{aligned}$$

Roughly speaking, **bias** measures how far off the target we hit on the average while **variance** measures how much fluctuation our estimator may show from

one sample to another.

4.2 Unbiased Estimators

In almost all real applications, the class of possible estimators for an **ESTIMANAL** is huge and the best estimator, i.e. the one that minimizes MSE no matter what the value of the **ESTIMANAL** is, almost never exists. Thus we try to reduce the class of potential estimators by improving a plausible restriction, for example $\text{Bias}(\hat{\Theta}_n) = 0$.

Definition. An estimator $\hat{\Theta}_n$ of an **ESTIMANAL** Θ is said to be **unbiased** if $E(\hat{\Theta}_n) = \Theta$, for all possible values of Θ .

Example 4.1. $X_i \stackrel{iid}{\sim} N(\mu, \sigma^2) \quad i = 1, 2, \dots, n$

Suppose both μ and σ^2 are unknown. Consider $\bar{X}_n = \frac{1}{n} \sum_{i=1}^n X_i$.

$$E(\bar{X}_n) = E\left(\frac{1}{n} \sum_{i=1}^n X_i\right) = \frac{1}{n} \sum_{i=1}^n \overbrace{E(X_i)}^{\mu} = \frac{1}{n} \cdot n\mu = \mu$$

Thus \bar{X}_n is an unbiased estimator of μ . As for the $\text{MSE}(\bar{X}_n)$, we need to find $\text{Var}(\bar{X}_n)$.

$$\begin{aligned} \text{Var}(\bar{X}_n) &= \text{Var}\left(\frac{1}{n} \sum_{i=1}^n X_i\right) = \frac{1}{n^2} \text{Var}\left(\sum_{i=1}^n X_i\right) \\ &= \frac{1}{n^2} \left[\sum_{i=1}^n \text{Var}(X_i) + 2 \cdot \sum_{1 \leq i < j \leq n} \overbrace{\text{Cov}(X_i, X_j)}^0 \right] && \text{Theorem 5.12(b) - page 271} \\ &= \frac{1}{n^2} \sum_{i=1}^n \text{Var}(X_i) && \prod_{i=1}^n X_i \\ &= \frac{1}{n^2} \sum_{i=1}^n \sigma^2 = \frac{1}{n^2} \cdot n\sigma^2 = \frac{\sigma^2}{n} && \text{identically distributed} \\ \implies \text{MSE}(\bar{X}_n) &= \text{Var}(\bar{X}_n) + \overbrace{\text{Bias}^2(\bar{X}_n)}^0 = \text{Var}(\bar{X}_n) = \frac{\sigma^2}{n} \end{aligned}$$

An inspection of the above calculation shows that for unbiased μ we only

require a common mean μ while for calculating the variance we would only require a common variance σ^2 and orthogonality, i.e:

$$\text{Cov}(X_i, X_j) = 0 \quad \text{where } i \neq j$$

Suppose X_1, \dots, X_n have the same mean value μ . Then:

$$E(\bar{X}_n) = E\left(\frac{1}{n} \sum_{i=1}^n X_i\right) = \frac{1}{n} \sum_{i=1}^n E(X_i) = \frac{1}{n} n \mu = \mu$$

Suppose further that X_1, \dots, X_n have the same variance σ^2 and $\text{Cov}(X_i, X_j) = 0$, $i \neq j$.

Then:

$$\begin{aligned} \text{Var}(\bar{X}_n) &= \text{Var}\left(\frac{1}{n} \sum_{i=1}^n X_i\right) = \frac{1}{n^2} \text{Var}\left(\sum_{i=1}^n X_i\right) \\ &= \frac{1}{n^2} \left[\sum_{i=1}^n \text{Var}(X_i) + 2 \sum_{1 \leq i < j \leq n} \text{Cov}(X_i, X_j) \right] && \text{Theorem 5.12(b) - Page 271} \\ &= \frac{1}{n^2} \sum_{i=1}^n \text{Var}(X_i) && \text{Orthogonality: i.e. } \text{Cov}(X_i, X_j) = 0 \text{ if } i \neq j \\ &= \frac{1}{n^2} \sum_{i=1}^n \sigma^2 = \frac{1}{n^2} n \sigma^2 = \frac{\sigma^2}{n} && \text{having the same variance} \\ &\implies \text{MSE}(\bar{X}_n) = \text{Var}(\bar{X}_n) = \frac{\sigma^2}{n} \end{aligned}$$

If X_1, \dots, X_n have the same mean value and variance and they are orthogonal.

4.3 Stein's Paradox

We will learn later that if $X_i \stackrel{iid}{\sim} N(\mu, \sigma^2)$ then $\bar{X}_n = \frac{1}{n} \sum_{i=1}^n X_i$ has many optimal properties. A paradox due to Charles Stein, however, shows that such a nice optimal properties are not preserved in higher dimensions. In fact if:

$$X_i \stackrel{iid}{\sim} N(\mu_x, 1), \quad Y_i \stackrel{iid}{\sim} N(\mu_y, 1) \text{ and } Z_i \stackrel{iid}{\sim} (\mu_z, 1)$$

then, we can find the biased estimators of $\begin{pmatrix} \mu_x \\ \mu_y \\ \mu_z \end{pmatrix}$ which are closer to $\begin{pmatrix} \mu_x \\ \mu_y \\ \mu_z \end{pmatrix}$ than $\begin{pmatrix} \bar{X}_n \\ \bar{Y}_n \\ \bar{Z}_n \end{pmatrix}$ for any $\begin{pmatrix} \mu_x \\ \mu_y \\ \mu_z \end{pmatrix}$. We may then say that $\begin{pmatrix} \bar{X}_n \\ \bar{Y}_n \\ \bar{Z}_n \end{pmatrix}$ is an **inadmissible estimator** of $\begin{pmatrix} \mu_x \\ \mu_y \\ \mu_z \end{pmatrix}$.

4.4 Admissibility

An estimator $\hat{\Theta}$ is called admissible if there is no estimator $\tilde{\Theta}$ such that:

$$MSE(\tilde{\Theta}) \leq MSE(\hat{\Theta}) \quad \text{for all possible values of } \Theta$$

and this inequality is strict for some values of Θ .

What this example tells us is that by allowing a bit of bias we may be able to reduce variance considerably and hence find an estimator which is closer to the target than the most natural unbiased estimator. Note that this phenomena happens only when the dimension is at least 3.

5 Lecture 4

We now want to restrict the class of estimators even further. Suppose X_1, \dots, X_n have the same mean μ and variance σ^2 and they are orthogonal; i.e. $Cov(X_i, X_j) = 0$, $i \neq j$. Consider $\tilde{X}_{n,\tilde{C}} = \sum_{i=1}^n C_i X_i$ and

$$\mathcal{C} = \left\{ \tilde{X}_{n,\tilde{C}} : \tilde{C} = (C_1, \dots, C_n) \in \mathbf{R}^n, \sum_{i=1}^n C_i = 1 \right\}$$

Note that

$$\begin{aligned} E(\tilde{X}_{n,\tilde{C}}) &= E\left(\sum_{i=1}^n C_i X_i\right) = \sum_{i=1}^n C_i E(X_i) \\ &= \sum_{i=1}^n C_i \mu = \mu \sum_{i=1}^n C_i = \mu \cdot 1 = \mu \end{aligned}$$

Thus $\tilde{X}_{n,\tilde{C}}$ is an unbiased estimator of μ for any $\tilde{C} \in \mathbf{R}^n$ as long as $\sum_{i=1}^n C_i = 1$. Then \mathcal{C} is the class of all unbiased linear estimators of μ . We want to find the best estimator with \mathcal{C} ; i.e.:

$$\underset{\tilde{C} \in \mathbf{R}^n}{\text{Min}} \text{MSE}(\tilde{X}_{n,\tilde{C}}) \quad \text{s.t.} \quad \sum_{i=1}^n C_i = 1 \quad (*)$$

First we note that $MSE(\tilde{X}_{n,\underline{C}}) = Var(\tilde{X}_{n,\underline{C}})$ since $\tilde{X}_{n,\underline{C}}$ is an unbiased estimator of μ when $\sum_{i=1}^n C_i = 1$. On the other hand:

$$\begin{aligned}
 Var(\tilde{X}_{n,\underline{C}}) &= Var\left(\sum_{i=1}^n C_i X_i\right) \\
 &= \sum_{i=1}^n C_i^2 Var(X_i) + 2 \sum_{1 \leq i < j \leq n} Cov(C_i X_i, C_j X_j) \quad \text{Theorem 5.12 page 271} \\
 &= \sum_{i=1}^n C_i^2 \sigma^2 + 2 \sum_{1 \leq i < j \leq n} \overbrace{C_i C_j Cov(X_i, X_j)}^0 \\
 &= \sigma^2 \sum_{i=1}^n C_i^2
 \end{aligned}$$

Thus (*) is equivalent to :

$$\underset{\underline{C} \in \mathbb{R}^n}{\text{Min}} \sigma^2 \sum_{i=1}^n C_i^2 \quad (**)$$

Using the *Lagrange Theorem*, (**) is equivalent to:

$$\underline{C} = (C_1, \dots, C_n) \in \mathbb{R}^n \quad \overbrace{\left\{ \sigma^2 \sum_{i=1}^n C_i + \lambda \left(\sum_{i=1}^n C_i - 1 \right) \right\}}^{\mathcal{C}_\lambda(\underline{C})}.$$

Note that:
$$\frac{\partial \mathcal{C}_\lambda(\underline{C})}{\partial C_i} = 2 \sigma^2 C_i + \lambda \quad , \quad i = 1, 2, 3, \dots$$

$$\frac{\partial}{\partial \lambda} \mathcal{C}_\lambda(\underline{C}) = \sum_{i=1}^n C_i - 1$$

$$\begin{cases} \frac{\partial}{\partial C_i} \mathcal{C}_\lambda(\underline{C}) = 2 \sigma^2 C_i + \lambda = 0 \quad , \quad i = 1, 2, 3, \dots \\ \frac{\partial}{\partial \lambda} \mathcal{C}_\lambda = 0 \implies \sum_{i=1}^n C_i = 1 \end{cases}$$

Thus $C_i = -\frac{\lambda}{2 \sigma^2}$, $i = 1, 2, 3, \dots, n$ and using the last equation:

$$\sum_{i=1}^n -\frac{\lambda}{2 \sigma^2} = 1 \implies \lambda = -\frac{2 \sigma^2}{n}$$

and therefore:

$$C_i = -\frac{\lambda}{2\sigma^2} = -\frac{-\frac{2\sigma^2}{n}}{2\sigma^2} = \frac{1}{n} \quad , \quad i = 1, 2, 3, \dots, n$$

We can further find:

$$\mathcal{H} = \left[\frac{\partial^2}{\partial C_i \partial C_j} \mathcal{C}_\lambda(\tilde{C}) \right] \quad , \quad i, j = 1, 2, \dots, n$$

and show that:

$$\begin{aligned} \tilde{x}^T \mathcal{H} \tilde{x} &\geq 0 \quad \forall \tilde{x} \in \mathbb{R}^n \\ &= 0 \quad \text{if and only if } \tilde{x} = 0 \end{aligned}$$

This then guarantees that $\tilde{C}^* = (\frac{1}{n}, \frac{1}{n}, \dots, \frac{1}{n})$ is indeed a minimizer; in fact, the *unique minimizer*. To summarize:

$$\tilde{X}_{n, \tilde{C}^*} = \sum_i i = 1^n \frac{1}{n} X_i = \frac{1}{n} \sum_{i=1}^n X_i = \bar{X}_n$$

Thus \bar{X}_n is the best unbiased linear estimator.

5.1 Estimating Variance

So far we confirmed ourselves to estimation of the population mean.

Now suppose we are interested in estimating variance from X_1, \dots, X_n where X_i s have the same mean value μ , the same variance σ^2 and they are orthogonal, i.e. $\text{Cov}(X_i, X_j) = 0$, $i \neq j$, then a *natural estimator* of:

$$\sigma^2 = \text{Var}(X) = \mathbb{E}[(x - \mu)^2]$$

is its sample counterpart, i.e.

$$S_{n,*}^2 = \frac{1}{n} \sum_{i=1}^n (X_i - \bar{X}_n)^2$$

Now the first question is if $S_{n,*}^2$ is an unbiased estimator of σ^2 , i.e. $\mathbb{E}(S_{n,*}^2) = \sigma^2$

$$\begin{aligned} (X_i - \mu)^2 &= \left[(X_i - \bar{X}_n) + (\bar{X}_n - \mu) \right]^2 \\ &= (X_i - \bar{X}_n)^2 + (\bar{X}_n - \mu)^2 + 2 \cdot (X_i - \bar{X}_n)(\bar{X}_n - \mu) \\ \sum_{i=1}^n (X_i - \mu)^2 &= \sum_{i=1}^n (X_i - \bar{X}_n)^2 + n(\bar{X}_n - \mu)^2 + 2 \cdot (\bar{X}_n - \mu) \overbrace{\sum_{i=1}^n (X_i - \bar{X}_n)}^0 \\ &= \sum_{i=1}^n (X_i - \bar{X}_n)^2 + n(\bar{X}_n - \mu)^2 \end{aligned} \quad (I)$$

Taking estimation we find:

$$\begin{aligned} \mathbb{E} \left[\sum_{i=1}^n (X_i - \mu)^2 \right] &= \mathbb{E}[n \cdot S_{n,*}^2] + \mathbb{E}[n(\bar{X}_n - \mu)^2] \\ RHS &= \sum_{i=1}^n \overbrace{\mathbb{E}(X_i - \mu)^2}^{\sigma^2} = n \cdot \sigma^2 \end{aligned} \quad (II)$$

Note that $\mathbb{E}(\bar{X}_n - \mu) = 0$, i.e. $\mathbb{E}(\bar{X}_n) = \mu$. Thus:

$$\mathbb{E}[n(\bar{X}_n - \mu)^2] = n \mathbb{E}[(\bar{X}_n - \mu)^2] = n \text{Var}(\bar{X}_n).$$

On the other hand $\text{Var}(\bar{X}_n) = \frac{\sigma^2}{n}$. We therefore have:

$$\mathbb{E}[n(\bar{X}_n - \mu)^2] = n \cdot \text{Var}(\bar{X}_n) = n \cdot \frac{\sigma^2}{n} = \sigma^2$$

and hence from (II):

$$n\sigma^2 = \mathbb{E}(n S_{n,*}^2) + \sigma^2$$

which implies:

$$\implies E(S_{n,*}^2) = \left(\frac{n-1}{n}\right)\sigma^2 = \left(1 - \frac{1}{n}\right)\sigma^2$$

meaning that $S_{n,*}^2$ is **NOT** an unbiased estimator of σ^2 .

Multiplying both sides of the last equation by the reciprocal of $(1 - \frac{1}{n})$ we find

$E(\frac{n}{n-1}S_{n,*}^2) = \sigma^2$. Note however that:

$$\frac{n}{n-1}S_{n,*}^2 = \frac{\cancel{n}}{n-1} \cdot \frac{1}{\cancel{n}} \sum_{i=1}^n (X_i - \bar{X}_n)^2 = \frac{1}{n-1} \sum_{i=1}^n (X_i - \bar{X}_n)^2$$

Thus $\boxed{S_n^2 = \frac{1}{n-1} \sum_{i=1}^n (X_i - \bar{X}_n)^2}$ is an unbiased estimator.

Question: Why $(n-1)$?

" $n-1$ " is the dimension of $\overbrace{\text{span}\{X_i - \bar{X}_n : i = 1, 2, \dots, n\}}^{\mathbb{W}}$.

$n-1 = \dim(\text{span } V)$. Note however $\dim(\text{span } W) = n$ where $W = X_i - \mu$, $i = 1, 2, \dots, n$.

We discuss these issues further in Chapter 11 where we learn about the regression.

So far we only considered sampling from one population. We may have samples from two or more populations and may want to make inference about differences between the populations.

Example 5.1.

Suppose we want to study the differences between the average salaries of men and women:

Men	Women
X_1	Y_1
\vdots	\vdots
X_m	Y_n

where X_i s have the common mean μ_x and Y_j s have the common mean μ_y .

We want to estimate $\mu_x - \mu_y$. The natural estimator is $\bar{X}_m - \bar{Y}_n$. Show that:

$$E[\bar{X}_m - \bar{Y}_n] = \mu_x - \mu_y$$

Hence $\bar{X}_m - \bar{Y}_n$ is an unbiased estimator of $\mu_x - \mu_y$.

Assume further that Xs and Ys are independent and Xs have common variance σ_x^2 and Ys have common variance σ_y^2 and $Cov(X_i, X_j) = 0$, $i \neq j$ and $Cov(Y_i, Y_j) = 0$, $i \neq j$.

Find $Var(\bar{X}_m - \bar{Y}_n)$. Hint: use *Thm* 5.12.

The difference between two proportions can be treated similarly. Note that proportions are essentially means of binary variables.

6 Lecture 5

7 Lecture 6

8 Lecture 7

9 Lecture 8

10 Lecture 9

11 Lecture 10

12 Lecture 11

13 Lecture 12